

#Bank Churn Model Using ML

# Import Library

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

#Import Data

```
df=pd.read_csv('https://raw.githubusercontent.com/YBI-Foundation/Dataset/main/Bank%20Churn%20Modelling.csv')
```

#Analyse Data

```
df.head()
```

	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	Num Of Products	Has Credit Card	Is Active Member
0	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1
1	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1
2	15619304	Onio	502	France	Female	42	8	159660.80	3	1	1
3	15701354	Boni	699	France	Female	39	1	0.00	2	0	1

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CustomerId            10000 non-null  int64
1   Surname               10000 non-null  object
2   CreditScore           10000 non-null  int64
3   Geography             10000 non-null  object
4   Gender               10000 non-null  object
5   Age                  10000 non-null  int64
6   Tenure               10000 non-null  int64
7   Balance              10000 non-null  float64
8   Num Of Products      10000 non-null  int64
9   Has Credit Card      10000 non-null  int64
10  Is Active Member     10000 non-null  int64
11  Estimated Salary     10000 non-null  float64
12  Churn               10000 non-null  int64
dtypes: float64(2), int64(8), object(3)
memory usage: 1015.8+ KB
```

```
df.duplicated('CustomerId').sum()
```

0

```
df=df.set_index('CustomerId')
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10000 entries, 15634602 to 15628319
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Surname               10000 non-null  object
1   CreditScore           10000 non-null  int64
2   Geography             10000 non-null  object
3   Gender               10000 non-null  object
4   Age                  10000 non-null  int64
5   Tenure               10000 non-null  int64
6   Balance              10000 non-null  float64
7   Num Of Products      10000 non-null  int64
```

```

8   Has Credit Card    10000 non-null int64
9   Is Active Member   10000 non-null int64
10  Estimated Salary    10000 non-null float64
11  Churn               10000 non-null int64
dtypes: float64(2), int64(7), object(3)
memory usage: 1015.6+ KB

```

```
# Encoding
```

```
df['Geography'].value_counts()
```

```

France    5014
Germany   2509
Spain     2477
Name: Geography, dtype: int64

```

```
df.replace({'Geography':{'France':2,'Germany':1,'Spain':0}},inplace=True)
```

```
df['Gender'].value_counts()
```

```

Male      5457
Female    4543
Name: Gender, dtype: int64

```

```
df.replace({'Gender':{'Male':0,'Female':1}},inplace=True)
```

```
df['Num Of Products'].value_counts()
```

```

1     5084
2     4590
3       266
4        60
Name: Num Of Products, dtype: int64

```

```
df.replace({'Num Of Products':{'1':0,2:1,3:1,4:1}},inplace=True)
```

```
df['Has Credit Card'].value_counts()
```

```

1     7055
0     2945
Name: Has Credit Card, dtype: int64

```

```
df['Is Active Member'].value_counts()
```

```

1     5151
0     4849
Name: Is Active Member, dtype: int64

```

```
df.loc[(df['Balance']==0),'Churn'].value_counts()
```


```

0     3117
1       500
Name: Churn, dtype: int64

```

```
df['Zero Balance']=np.where(df['Balance']>0,1,0)
```

```
df['Zero Balance'].hist()
```



```
df.groupby(['Churn', 'Geography']).count()
```

		Surname	CreditScore	Gender	Age	Tenure	Balance	Num Of Products	Has Credit Card	Is Active Member	Estimated Salary
Churn	Geography										
0	0	2064	2064	2064	2064	2064	2064	2064	2064	2064	2064
	1	1695	1695	1695	1695	1695	1695	1695	1695	1695	1695
	2	4204	4204	4204	4204	4204	4204	4204	4204	4204	4204
1	0	413	413	413	413	413	413	413	413	413	413
	1	814	814	814	814	814	814	814	814	814	814

```
#Define Label And Features
df.columns

Index(['Surname', 'CreditScore', 'Geography', 'Gender', 'Age', 'Tenure',
      'Balance', 'Num Of Products', 'Has Credit Card', 'Is Active Member',
      'Estimated Salary', 'Churn', 'Zero Balance'],
      dtype='object')

X=df.drop(['Surname','Churn'],axis=1)

y=df['Churn']


X.shape,y.shape

((10000, 11), (10000,))

df['Churn'].value_counts()

0    7963
1    2037
Name: Churn, dtype: int64

sns.countplot(x = 'Churn',data=df);
```



```
#Random Under Sampling
from imblearn.under_sampling import RandomUnderSampler

rus = RandomUnderSampler(random_state=2529)

X_rus , y_rus = rus.fit_resample(X,y)

X_rus.shape,y_rus.shape,X.shape,y.shape

((4074, 11), (4074,)), (10000, 11), (10000,))

y.value_counts()


0    7963
1    2037
Name: Churn, dtype: int64

y_rus.value_counts()

0    2037
1    2037
Name: Churn, dtype: int64
```

```
y_rus.plot(kind='hist')
```

<Axes: ylabel='Frequency'>



```
#Random Over Sampling
from imblearn.over_sampling import RandomOverSampler

ros=RandomOverSampler(random_state=2529)

X_ros,y_ros=ros.fit_resample(X,y)

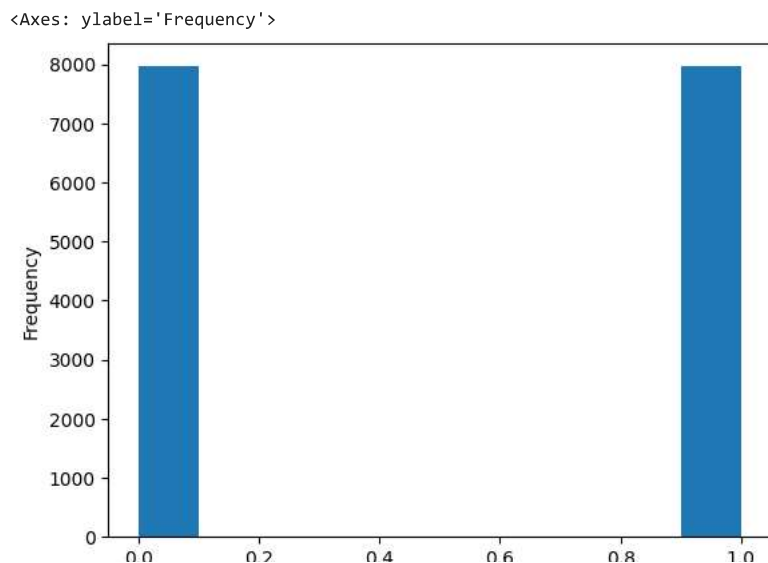
X_ros.shape,y_ros.shape,X.shape,y.shape

((15926, 11), (15926,)), (10000, 11), (10000,))

y.value_counts()

0    7963
1    2037
Name: Churn, dtype: int64

y_ros.plot(kind='hist')
```



```
#Train Test Split
from sklearn.model_selection import train_test_split

#Split Original Data
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=25)

#Split Random Under Sample Data
X_train_rus,X_test_rus,y_train_rus,y_test_rus=train_test_split(X_rus,y_rus,test_size=0.3)

#Split Random Over Sample Data
X_train_ros,X_test_ros,y_train_ros,y_test_ros=train_test_split(X_ros,y_ros,test_size=0.3)

#Standardize Features
from sklearn.preprocessing import StandardScaler

sc=StandardScaler()

#Standardize Original Data
X_train[['CreditScore','Age','Tenure','Balance','Estimated Salary']]=sc.fit_transform(X_train[['CreditScore','Age','Tenure','Balance','Estima

#Standardize Original Data
X_test[['CreditScore','Age','Tenure','Balance','Estimated Salary']]=sc.fit_transform(X_test[['CreditScore','Age','Tenure','Balance','Estimate

#Standardize Random Under Sample Data
X_train_rus[['CreditScore','Age','Tenure','Balance','Estimated Salary']]=sc.fit_transform(X_train_rus[['CreditScore','Age','Tenure','Balance'
X_test_rus[['CreditScore','Age','Tenure','Balance','Estimated Salary']]=sc.fit_transform(X_test_rus[['CreditScore','Age','Tenure','Balance','

#standardize Random over Sample Data
X_train_ros[['CreditScore','Age','Tenure','Balance','Estimated Salary']]=sc.fit_transform(X_train_ros[['CreditScore','Age','Tenure','Balance'
X_test_ros[['CreditScore','Age','Tenure','Balance','Estimated Salary']]=sc.fit_transform(X_test_ros[['CreditScore','Age','Tenure','Balance','

X_train_rus.head()
```

X\_train\_ros.head()

	CreditScore	Geography	Gender	Age	Tenure	Balance	Num Of Products	Has Credit Card	Is Active Member	Estimated Salary	
3292	-1.570296		2	1	-0.012695	-1.362155	1.542957	1	1	0	1.42697
1681	0.324677		0	1	-0.855625	-0.676186	-1.315602	1	1	1	-0.70382
12685	0.232489		2	0	1.017554	1.038738	1.445272	0	0	0	-0.76067
8656	0.109572		0	0	-0.949284	-1.362155	0.057618	0	1	1	0.65299

X\_test\_ros.head()

	CreditScore	Geography	Gender	Age	Tenure	Balance	Num Of Products	Has Credit Card	Is Active Member	Estimated Salary	Z
15904	749		1	0	54	3	144768.94	0	1	0	93336.30
7215	548		1	0	32	2	98986.28	0	1	1	55867.38
1704	654		2	1	29	4	96974.97	0	0	1	141404.07
12499	850		2	1	42	8	0.00	0	1	0	19632.64

X\_test\_rus.head()

	CreditScore	Geography	Gender	Age	Tenure	Balance	Num Of Products	Has Credit Card	Is Active Member	Estimated Salary	
1480	1.330783		0	0	-0.816414	1.418574	-1.283758	1	1	0	-1.566655
3012	0.434908		2	1	0.285452	1.418574	-1.283758	1	0	0	0.067755
2555	0.342231		1	1	0.928207	-1.419734	1.065661	0	0	0	0.185799
549	0.795318		2	0	-0.357303	-1.064946	-1.283758	0	0	0	0.383623

#Support Vector Machine Classifier

from sklearn.svm import SVC

svc=SVC()

svc.fit(X\_train,y\_train)

▼ SVC

SVC()

y\_pred = svc.predict(X\_test)

y\_pred=svc.predict(X\_test)

#Model Accuracy

from sklearn.metrics import confusion\_matrix,classification\_report

confusion\_matrix(y\_test,y\_pred)

array([[2374, 45],  
[ 421, 160]])

print(classification\_report(y\_test,y\_pred))

	precision	recall	f1-score	support
0	0.85	0.98	0.91	2419
1	0.78	0.28	0.41	581
accuracy			0.84	3000
macro avg	0.81	0.63	0.66	3000
weighted avg	0.84	0.84	0.81	3000

```
#Hyperparameter Tunning
```

```
from sklearn.model_selection import GridSearchCV
```

```
param_grid = {'C':[0.1,1,10],
              'gamma':[1,0.1,0.01],
              'kernel':['rbf'],
              'class_weight':['balanced']}
```

```
grid = GridSearchCV(SVC(),param_grid,refit=True,verbose=2,cv=2)
grid.fit(X_train,y_train)
```

```
Fitting 2 folds for each of 9 candidates, totalling 18 fits
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.5s
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.4s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.1s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.1s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.2s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.2s
[CV] END ...C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.3s
[CV] END ...C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 2.2s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.3s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.9s
[CV] END C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.0s
[CV] END ..C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.0s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.2s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 1.3s
[CV] END ..C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.0s
[CV] END ..C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 1.0s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.0s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 1.0s
```

```
GridSearchCV
  estimator: SVC
    SVC
```

```
print(grid.best_estimator_)
```

```
SVC(C=10, class_weight='balanced', gamma=1)
```

```
grid_predictions=grid.predict(X_test)
```

```
confusion_matrix(y_test,grid_predictions)
```

```
array([[2166, 253],
       [ 365, 216]])
```

```
print(classification_report(y_test,grid_predictions))
```

	precision	recall	f1-score	support
0	0.86	0.90	0.88	2419
1	0.46	0.37	0.41	581
accuracy			0.79	3000
macro avg	0.66	0.63	0.64	3000
weighted avg	0.78	0.79	0.79	3000

```
#Model with Random Under Sampling
```

```
svc_rus=SVC()
```

```
svc_rus.fit(X_train_rus,y_train_rus)
```

▼ SVC  
SVC()

```
y_pred_rus=svc_rus.predict(X_test_rus)
```

```
#Model Accuracy
```

```
confusion_matrix(y_test_rus,y_pred_rus)
```

```
array([[476, 150],
       [164, 433]])
```

```
print(classification_report(y_test_rus,y_pred_rus))
```

```

              precision    recall  f1-score   support

     0       0.74         0.76         0.75         626
     1       0.74         0.73         0.73         597

 accuracy          0.74
 macro avg         0.74
 weighted avg      0.74
```

```
#Hyperparameter Tunning
```

```
param_grid = {'C':[0.1,1,10],
              'gamma':[1,0.1,0.01],
              'kernel':['rbf'],
              'class_weight':['balanced']}
```

```
grid_rus = GridSearchCV(SVC(),param_grid,refit=True,verbose=2,cv=2)
grid_rus.fit(X_train_rus,y_train_rus)
```

```

Fitting 2 folds for each of 9 candidates, totalling 18 fits
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.4s
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.4s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.3s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.4s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.3s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s
[CV] END ...C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.2s
[CV] END ...C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.2s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END ..C=1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s
[CV] END .C=1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.2s
[CV] END ...C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 0.2s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END .C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 0.2s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s
[CV] END C=10, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 0.2s
```

► GridSearchCV  
► estimator: SVC  
    ► SVC

```
print(grid_rus.best_estimator_)
```

```
SVC(C=10, class_weight='balanced', gamma=0.01)
```

```
grid_predictions_rus =grid_rus.predict(X_test_rus)
```

```
confusion_matrix(y_test_rus,grid_predictions_rus)
```

```
array([[462, 164],
       [170, 427]])
```

```
print(classification_report(y_test_rus,grid_predictions_rus))
```



	precision	recall	f1-score	support
0	0.73	0.74	0.73	626
1	0.72	0.72	0.72	597
accuracy			0.73	1223
macro avg	0.73	0.73	0.73	1223
weighted avg	0.73	0.73	0.73	1223

```
#Model with Random Over Sampling
svc_ros=SVC()
```

```
svc_ros.fit(X_train_ros,y_train_ros)
```

▼ SVC

SVC()

```
y_pred_ros=svc_ros.predict(X_test_ros)
```

```
#Model Accuracy
confusion_matrix(y_test_ros,y_pred_ros)
```

```
array([[2375,  0],
       [2403,  0]])
```

```
print(classification_report(y_test_ros,y_pred_ros))
```

	precision	recall	f1-score	support
0	0.50	1.00	0.66	2375
1	0.00	0.00	0.00	2403
accuracy			0.50	4778
macro avg	0.25	0.50	0.33	4778
weighted avg	0.25	0.50	0.33	4778

```
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-de
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-de
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-de
_warn_prf(average, modifier, msg_start, len(result))
```

```
#Hyperparameter Tunning
```

```
#Hyperparameter Tunning
```

```
param_grid = {'C':[0.1,1,10],
              'gamma':[1,0.1,0.01],
              'kernel':['rbf'],
              'class_weight':['balanced']}
```

```
grid_ros = GridSearchCV(SVC(),param_grid,refit=True,verbose=2,cv=2)
grid_ros.fit(X_train_ros,y_train_ros)
```

```

Fitting 2 folds for each of 9 candidates, totalling 18 fits
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 3.8s
[CV] END ..C=0.1, class_weight=balanced, gamma=1, kernel=rbf; total time= 4.8s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 2.7s
[CV] END C=0.1, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 2.7s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 3.1s
[CV] END C=0.1, class_weight=balanced, gamma=0.01, kernel=rbf; total time= 4.2s
[CV] END ...C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 3.2s
[CV] END ...C=1, class_weight=balanced, gamma=1, kernel=rbf; total time= 3.2s

print(grid_ros.best_estimator_)

SVC(C=10, class_weight='balanced', gamma=1)
[CV] END C=10, class_weight=balanced, gamma=1, kernel=rbf; total time= 4.0s

grid_predictions_ros=grid_ros.predict(X_test_ros)
[CV] END C=10, class_weight=balanced, gamma=0.1, kernel=rbf; total time= 4.0s

confusion_matrix(y_test_ros,grid_predictions_ros)

array([[2375,  0],
       [2403,  0]])
      > SVC |

print(classification_report(y_test_ros,grid_predictions_ros))

              precision    recall  f1-score   support

     0         0.50         1.00         0.66         2375
     1         0.00         0.00         0.00         2403

 accuracy          0.50          0.50          0.50          4778
 macro avg         0.25         0.50         0.33          4778
 weighted avg         0.25         0.50         0.33          4778

/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-de
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-de
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-de
_warn_prf(average, modifier, msg_start, len(result))

```

#Lets Compare

```

print(classification_report(y_test,y_pred))

              precision    recall  f1-score   support

     0         0.85         0.98         0.91         2419
     1         0.78         0.28         0.41          581

 accuracy          0.84          0.84          0.84         3000
 macro avg         0.81         0.63         0.66         3000
 weighted avg         0.84         0.84         0.81         3000

print(classification_report(y_test,grid_predictions))

              precision    recall  f1-score   support

     0         0.86         0.90         0.88         2419
     1         0.46         0.37         0.41          581

 accuracy          0.79          0.79          0.79         3000
 macro avg         0.66         0.63         0.64         3000
 weighted avg         0.78         0.79         0.79         3000

print(classification_report(y_test_rus,y_pred_rus))

              precision    recall  f1-score   support

     0         0.74         0.76         0.75          626
     1         0.74         0.73         0.73          597

 accuracy          0.74          0.74          0.74         1223
 macro avg         0.74         0.74         0.74         1223
 weighted avg         0.74         0.74         0.74         1223

```

```
print(classification_report(y_test_rus,grid_predictions_rus))
```

	precision	recall	f1-score	support
0	0.73	0.74	0.73	626
1	0.72	0.72	0.72	597
accuracy			0.73	1223
macro avg	0.73	0.73	0.73	1223
weighted avg	0.73	0.73	0.73	1223

```
print(classification_report(y_test_ros,y_pred_ros))
```

	precision	recall	f1-score	support
0	0.50	1.00	0.66	2375
1	0.00	0.00	0.00	2403
accuracy			0.50	4778
macro avg	0.25	0.50	0.33	4778
weighted avg	0.25	0.50	0.33	4778

```
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-de
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-de
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-de
_warn_prf(average, modifier, msg_start, len(result))
```

```
print(classification_report(y_test_ros,grid_predictions_ros))
```

	precision	recall	f1-score	support
0	0.50	1.00	0.66	2375
1	0.00	0.00	0.00	2403
accuracy			0.50	4778
macro avg	0.25	0.50	0.33	4778
weighted avg	0.25	0.50	0.33	4778

```
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-de
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-de
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-de
_warn_prf(average, modifier, msg_start, len(result))
```